

Université Côte d'Azur

Faculty of Science and Technology
Department of Computer Science

Master 1 Project Report

AI Game Programming

16-Hole Mancala Game

Implementation of Advanced Search Algorithms

Prepared by:

BEN SALAH Mohamed Dhia
Mohamed Mehdi Khliaa
Master 1 Computer Science

Supervised by:

Prof. Regin Jean-Charles
Department of Computer Science

Contents

1	Introduction	1
1.1	Project Overview and Objectives	1
2	Game Rules and Mechanics	2
2.1	Board Configuration and Seed Types	2
2.2	Capturing and Victory Conditions	2
3	AI Algorithm: Min-Max with Alpha-Beta Pruning	4
3.1	Algorithm Choice and Justification	4
3.1.1	Algorithm Description	5
3.2	Evaluation Function Design	5
3.2.1	Function Design and Weight Justification	6
3.2.2	Alternative Components Considered	7
4	System Architecture	8
4.1	Design Decisions	8
4.2	Project Structure and Communication	8
5	Performance and Optimization	10
5.1	Performance Analysis and Optimization Techniques	10
6	Conclusion	11
6.1	Summary and Lessons Learned	11
6.2	Future Improvements	11

Chapter 1

Introduction

1.1 Project Overview and Objectives

This project implements a sophisticated **16-hole Mancala game** with an advanced artificial intelligence algorithm. The game is a variant of the traditional Mancala board game, featuring unique rules and mechanics that provide an interesting challenge for AI development.

The primary objective is to develop an efficient AI algorithm for game playing using the Min-Max approach with Alpha-Beta Pruning, a technique commonly used in adversarial game AI. The main goals are:

- Implement a complete 16-hole Mancala game engine
- Develop a Min-Max AI with Alpha-Beta Pruning for intelligent gameplay
- Create an arbitration system for bot-vs-bot matches
- Optimize the algorithm for real-time decision making (3-second timeout)
- Implement efficient evaluation functions and pruning strategies

Table 1.1: Technologies and Tools Used

Technology	Purpose	Details
C++17	Game Engine & AI	High-performance implementation
Java	Arbitration System	Process management & timing
STL Containers	Data Structures	Maps, vectors, queues
Chrono Library	Time Management	Timeout handling

Chapter 2

Game Rules and Mechanics

2.1 Board Configuration and Seed Types

The game board consists of **16 holes** numbered from 1 to 16, arranged in a circular pattern moving clockwise. Each player controls 8 holes:

Table 2.1: Hole Assignment per Player

Player	Holes	Description
Player 1	1, 3, 5, 7, 9, 11, 13, 15	Odd-numbered holes
Player 2	2, 4, 6, 8, 10, 12, 14, 16	Even-numbered holes

Each hole initially contains **6 seeds**: 2 of each color. The three seed colors have different sowing behaviors:

Table 2.2: Seed Colors and Their Behaviors

Color	Symbol	Sowing Behavior
Red	R	Distributed to ALL holes (clockwise)
Blue	B	Distributed to OPPONENT's holes only
Transparent	T	Played as Red (TR) or Blue (TB)

Moves are expressed using the format [Hole Number] [Color], for example: 14B (play Blue from hole 14), 3R (play Red from hole 3), 5TR (play Transparent as Red from hole 5).

2.2 Capturing and Victory Conditions

The capture rules are as follows:

1. A capture occurs when a sown seed brings a hole's total to **exactly 2 or 3 seeds**
2. Captures are **chained**: if the previous hole also has 2–3 seeds, those are captured too

3. Captures can occur in **any hole** (including own holes)
4. Seeds are captured **regardless of color**
5. **Starving the opponent is allowed**

Table 2.3: Game End Conditions

Condition	Result
Player captures ≥ 49 seeds	Immediate Victory
Board has < 10 seeds	Game ends, highest score wins
400 moves played	Game ends, highest score wins
Both players capture ≥ 48 seeds	Draw

Chapter 3

AI Algorithm: Min-Max with Alpha-Beta Pruning

3.1 Algorithm Choice and Justification

This project implements the **Min-Max algorithm with Alpha-Beta Pruning**, an adversarial search technique highly effective for two-player games. The choice was motivated by comparing alternatives:

Table 3.1: Algorithm Selection Justification

Algorithm	Limitation	Why Not Chosen
BFS / DFS	No adversarial modeling	Does not consider opponent's optimal responses
Pure Min-Max	Exponential complexity	Too slow for 3-second timeout constraint
MCTS	Requires many simulations	Needs more time to converge; less deterministic
Reinforcement Learning	Training data required	No pre-existing game database available

Key advantages of Alpha-Beta Pruning:

1. **Optimal Play Guarantee:** Ensures the AI plays optimally against a perfect opponent
2. **Deterministic Behavior:** Produces consistent, reproducible results
3. **Efficiency:** Reduces search space from $O(b^d)$ to $O(b^{d/2})$ in best case
4. **No Training Required:** Works immediately with a well-designed evaluation function

Why Depth 5? Through empirical testing, depth 5 provides optimal trade-off: sufficient lookahead for strategic play (1-2s avg) while consistently staying within the 3-second limit. Depth 6+ risks timeout in complex positions.

3.1.1 Algorithm Description

The Min-Max algorithm assumes optimal play from both players. Alpha-Beta pruning eliminates branches that cannot affect the final decision:

- **Maximizer:** Tries to maximize the evaluation score
- **Minimizer:** Tries to minimize the evaluation score
- **Alpha (α):** Best value for maximizer; **Beta (β):** Best value for minimizer
- **Cutoff:** When $\beta \leq \alpha$, prune remaining branches

Algorithm 1 Alpha-Beta Pruning

```

1: function ALPHABETA(state, depth, α, β, maximizing)
2:   if depth = 0 or isTerminal(state) then
3:     return evaluate(state)
4:   end if
5:   if maximizing then
6:     value  $\leftarrow -\infty$ 
7:     for each move m in getMoves(state) do
8:       child  $\leftarrow$  apply(state, m)
9:       value  $\leftarrow$  max(value, ALPHABETA(child, depth - 1, α, β, false))
10:      α  $\leftarrow$  max(α, value)
11:      if β  $\leq \alpha$  then
12:        break                                ▷ Beta cutoff
13:      end if
14:    end for
15:    return value
16:  else
17:    value  $\leftarrow +\infty$ 
18:    for each move m in getMoves(state) do
19:      child  $\leftarrow$  apply(state, m)
20:      value  $\leftarrow$  min(value, ALPHABETA(child, depth - 1, α, β, true))
21:      β  $\leftarrow$  min(β, value)
22:      if β  $\leq \alpha$  then
23:        break                                ▷ Alpha cutoff
24:      end if
25:    end for
26:    return value
27:  end if
28: end function

```

3.2 Evaluation Function Design

The evaluation function is the **most critical component** of the Min-Max algorithm. It determines how the AI perceives the value of a game state and directly impacts the

quality of decisions. A poorly designed function leads to missed opportunities, strategic blindness, and suboptimal play.

3.2.1 Function Design and Weight Justification

The evaluation function was designed after careful analysis of the game mechanics:

$$\text{Score}(s, p) = \underbrace{10 \times (C_p - C_o)}_{\text{Capture Advantage}} + \underbrace{2 \times (S_p - S_o)}_{\text{Board Control}} \quad (3.1)$$

where C_p = seeds captured by player, C_o = seeds captured by opponent, S_p = seeds on player's holes, S_o = seeds on opponent's holes.

Table 3.2: Evaluation Function Weight Justification

Component	Weight	Justification
Captured Seeds	10	Primary objective. Captures are permanent and directly determine the winner. A player needs 49 seeds to win immediately.
Board Control	2	Secondary factor. Seeds on your holes provide future move options but can be captured. Less permanent than captures.

Why the 5:1 ratio (10:2)? Through experimentation:

- **Equal weights (1:1):** AI prioritizes accumulating seeds rather than capturing → passive play
- **Capture only (10:0):** AI ignores board position → misses setups for future captures
- **Ratio 5:1 (10:2):** Optimal balance—prioritizes captures while maintaining strategic presence
- **Higher ratios (20:1):** Too aggressive → ignores defensive positioning

3.2.2 Alternative Components Considered

Table 3.3: Evaluation Components Analysis

Component	Included?	Reasoning
Captured Seeds Diff	Yes	Core winning condition
Board Seed Diff	Yes	Indicates move options and potential
Hole Distribution	No	Computation overhead; marginal benefit
Capture Threats	No	Increases complexity; depth handles this
Mobility (# of moves)	No	All holes usually have moves; not discriminating

For terminal states: Victory returns $+\infty$, Defeat returns $-\infty$, Draw returns 0. This ensures the algorithm always prefers winning moves and avoids losing moves regardless of other factors.

Chapter 4

System Architecture

4.1 Design Decisions

Why C++ for the Game Engine? C++ was chosen for maximum performance given the 3-second timeout constraint:

Table 4.1: Language Choice Justification

Language	Advantage	Limitation
C++ (Chosen)	Maximum performance, low-level control	More complex memory management
Python	Rapid development, readable	Too slow for deep search in 3s
Java	Good performance, portable	Slower than C++, JVM overhead

Why std::map for Game State? Maps provide cleaner code for sparse access patterns (holes 1-16 with 3 colors each). The slight performance cost vs arrays is offset by significantly improved maintainability.

4.2 Project Structure and Communication

Table 4.2: Project File Structure

File	Description
config.h	Game configuration and AI parameters
game_rules.h	Game rules and GameState class
game_engine.h	Move execution and capture logic
ai_algorithms.h	All AI algorithm implementations
bot.cpp	Standalone bot for external arbitration
Arbitre.java	Java referee for bot-vs-bot matches

The Java `Arbitre` class manages bot-vs-bot matches with 3-second timeout per move, automatic disqualification for invalid moves, and detailed game logging. Communication uses simple text protocol: `START` (game begins), move strings like `14B`, `NOMOVE` (no valid moves), and `RESULT` (game end).

Chapter 5

Performance and Optimization

5.1 Performance Analysis and Optimization Techniques

The Alpha-Beta pruning optimization significantly reduces the number of nodes evaluated, allowing deeper search within time constraints. Each optimization was chosen based on its impact-to-complexity ratio:

Table 5.1: Optimization Techniques with Justification

Technique	Implementation	Why This Approach?
Alpha-Beta Pruning	Reduces nodes by 50%	Essential for reaching depth 5 in time
Pre-computed Arrays	Static hole arrays	Avoids repeated computation of player holes
Inlined Evaluation	Critical path optimization	Eliminates function call overhead
Check Every 500 Nodes	Timeout monitoring	Balance between safety and overhead

Time Management: Each bot has a strict 3-second timeout. Exceeding this limit results in immediate disqualification. We experimented with timeout check intervals:

- Every node: 15% slowdown (too much overhead)
- Every 100 nodes: 5% slowdown
- **Every 500 nodes: <1% overhead** (optimal choice)
- Every 1000+ nodes: risk of exceeding timeout

Chapter 6

Conclusion

6.1 Summary and Lessons Learned

This project successfully implements a complete 16-hole Mancala game with an advanced Min-Max algorithm using Alpha-Beta Pruning. The key achievements include a complete game engine, efficient AI with Alpha-Beta optimization, robust arbitration system, and optimized real-time performance.

Key Lessons:

1. **Evaluation Function is King:** A mediocre algorithm with an excellent evaluation function outperforms a sophisticated algorithm with a poor one. We spent 60% of development time tuning the evaluation weights.
2. **Alpha-Beta is Non-Negotiable:** Without pruning, depth 5 would require exploring ~ 10 million nodes. With pruning, this drops to $\sim 3,000$ nodes—a 3000x improvement.
3. **Simplicity Wins:** We initially tried complex evaluation features (mobility, threat detection). The simpler 2-component function performed better due to faster computation allowing deeper search.
4. **Time Management is Critical:** A single timeout means disqualification. We learned to be conservative with depth and aggressive with pruning.

6.2 Future Improvements

Table 6.1: Potential Future Enhancements

Enhancement	Description
Transposition Tables	Cache evaluated positions for reuse
Move Ordering	Evaluate promising moves first
Machine Learning	Train evaluation function with game data
Monte Carlo Tree Search	Add MCTS for comparison

Thank you for reading